

Leadership workshop, one day

Design an AI-Native Software Factory for Your Organization

A working session for people who decide how engineering works. They do not write code here. They leave with an operating model, an autonomy matrix and a sequenced roadmap.

0:00 to 1:15

What you already have, unwritten

Goal Each team's current rules are on paper, and the variation between teams is visible to everyone in the room.

Say

- Open by asserting that they already have an operating model. It is informal, it differs by team, and nobody has compared the versions.
- Have each team write theirs on a card: what a machine may do unattended, what needs approval, and who grants exceptions.
- Put the cards side by side and let the spread speak. That gap is the reason for the day, and it is more persuasive than any slide.

Run Nothing technical. Cards on a wall.

Room does

- Produces one card per team and reads two others aloud.
- Marks every place two teams contradict each other.

Ask *Where two of your teams disagree, which one is right, and who would decide that today?*

1:30 to 3:00

The autonomy matrix

Goal Their estate is sorted by consequence, with approval placed where the risk actually is.

- Say**
- Kill the difficulty axis early, using an example from their own estate rather than a generic one.
 - Run the sort as a group exercise across their real services, and let the arguments happen. The arguments are the product.
 - Push back whenever everything lands in the strictest column. A model where nothing may proceed is a model nobody will follow.
-

Run

```
$ node steps/01-boundary/demo.ts graph
```

Show the worked example once, as a shape to copy, then put it away.

- Room does**
- Places every significant service in a column, with a named owner.
 - Agrees who may move something between columns and how that gets recorded.
-

Ask *Which service caused the longest argument? That is the one your written policy most needs to cover.*

3:45 to 5:00

Protecting the standard you already have

Goal	They know which checks must run on every change and which belong before release, and where their current setup fails open.
Say	• Ask what a machine could produce today that would pass their pipeline and still be wrong. Let the room be uncomfortable. • Introduce failing closed as a governance property rather than a technical one, because a missing check is invisible in every report they receive. • Explain the independent read in terms they own: an assurance step, with a cost per change they can decide to pay.
Run	<pre>\$ bash steps/04-verification/gates.sh --without-tests</pre>
Room does	• Sees a missing check produce something that is neither a pass nor a failure. • Lists the checks that will be required across every team.
Ask	<i>How would you learn, this month, that one of your teams quietly disabled a required check?</i>

5:00 to 5:30

Who is accountable for a change nobody wrote

Goal Accountability is assigned explicitly, and the capacity limit is a number rather than a feeling.

Say

- Put the question bluntly, because it is the one their board will ask them. A machine wrote it, someone approved it, and something went wrong.
- Settle it the same way as any other change: a named engineer approved it and owns it.
- Then give them the constraint that follows. Approval capacity is finite, so the pipeline must stop producing when it is full.

Run

```
$ node steps/01-boundary/demo.ts capacity
```

Room does

- Writes the accountability rule in one sentence.
- Sets a review capacity per team and says how it will be measured.

Ask *How many changes per week can each of your teams genuinely stand behind?*

5:30 to 6:00

Sequencing it

Goal	They leave with an ordered plan whose first step is small enough to happen next week.
-------------	---

- | | |
|------------|--|
| Say | <ul style="list-style-type: none">• Insist the first move is boring and cheap. Trust is built by a small thing that works, not by a programme.• Have them pick the first workflow by two criteria: it happens often, and getting it wrong is recoverable.• Name what they will measure before starting, so the second conversation is about numbers rather than impressions. |
|------------|--|
-

Run	Nothing. The roadmap template on screen.
------------	--

- | | |
|------------------|--|
| Room does | <ul style="list-style-type: none">• Fills in the first three steps with owners and dates.• Agrees the baseline they will compare against later. |
|------------------|--|
-

Ask	<i>What will you be able to show in six weeks, and who needs to see it?</i>
------------	---

2:20 to 2:40


Four shapes, one afternoon

Goal Each leader has chosen a shape for one real team and can defend the cost it imposes.

- Say**
- Switch the shape live and re-run the same item so the room sees the rules doing the work rather than the model.
 - None of the four is correct. Each is a position about who carries the cost of being wrong.
 - Push back on anyone choosing the strictest by default: it has a price, and it is paid in throughput by people who are not in this room.

Run

```
$ make profile NAME=solo
$ make profile NAME=enterprise
```



The same item, under the strictest of the four shapes.

Show on the board Settings to the two switches to Work to the money item

A shape is a commit, not a setting

make profile writes the charter and the ownership graph and pushes them, because a run clones from the remote. A shape that exists only on somebody's laptop is a shape no agent will ever read.

The queue cap is the number to argue about

Every shape declares how many decisions may wait on a person. That number, not the model, is what decides whether the factory helps or buries the team.

Nobody in the room pays the cost of the strictest shape

Pushing everything to refuse feels responsible and is paid for in throughput by people who are not at this table. Ask who they are and whether they agreed.

Room does

- Picks a shape for one named team and writes the trade-off it accepts.
 - Names the first rule they would relax after a quarter of evidence.
-

Ask

Who in your organisation currently pays for a change being wrong, and do they know?

Concepts: the six jobs, a written policy, grading autonomy, paths off limits, capacity, accountability, measuring, ownership.

Evidence to bring: what adoption measured, felt against measured speed, the named risk classes, routing as published practice.